



Pandas Cheat Sheet

From Basics to Advanced

Practical mini-tables + code + results + pro tips

Created by
Naveen Dhawan



DataFrame Basics

```
Pandas

import pandas as pd

data = {
    "Name": ["Alice", "Bob", "Charlie"],
    "Age": [28, 35, 22],
    "Country": ["India", "USA", "UK"],
}

df = pd.DataFrame(data)

print(df.shape)
print(df.dtypes)
print(df.head())
```

Result (snippet)

```
(3, 3)
Name      object
Age        int64
Country    object
dtype: object
```

	Name	Age	Country
0	Alice	28	India
1	Bob	35	USA
2	Charlie	22	UK

Explanation

- **df.shape** shows the number of rows × columns.
- **df.dtypes** shows column types.
- Use **df.head()** and **df.tail()** to quickly inspect data.

Best Practice: 💡

Always print a small preview of **df.head()** before deep analysis.

Rename columns & fix dtypes

```
Pandas

df = pd.DataFrame({
    'full_name': ['Alice Johnson', 'Bob Smith', 'Charlie Brown'],
    'age_str': ['28', '35', '22'],
    'salary_str': ['50000', '72000', '31000']
})
# Rename
df.rename(columns={'full_name': 'Name', 'age_str': 'Age',
                  'salary_str': 'Salary'}, inplace=True)
# Convert types
df['Age'] = pd.to_numeric(df['Age'])
df['Salary'] = df['Salary'].astype(int)
print(df)
```

Naveen Dhawan

Result

	Name	Age	Salary
0	Alice Johnson	28	50000
1	Bob Smith	35	72000
2	Charlie Brown	22	31000

- **df.dtypes** → Name: object, Age: int64, Salary: int64

Explanation/Tip

- Use `errors='coerce'` to identify bad values (will become NaN).
- Use **`.astype()`** to correct incorrect data types before joining.
- Common error ⚠️: “String money values – remove commas before `to_numeric`.”



Indexing (set_index, reset_index, reindex)

```
Pandas

df = pd.DataFrame({'OrderID':[1001,1002,1003],
                  'Product':
                  ['Phone','Laptop','Tablet'],
                  'Qty':[2,1,3]})

# set index
df.set_index('OrderID', inplace=True)
# reset index
df_reset = df.reset_index()
```

Result (snippet)

OrderID	Product	Qty
1001	Phone	2
1002	Laptop	1
1003	Tablet	3

- After **set_index**: index values = [1001,1002,1003];
columns = Product, Qty

When to use:

Use **set_index** for natural keys (OrderID) and reset before concatenating disparate sources.

Explanation/Tip

- **reset_index(drop=True)** removes the index and creates a default integer index.
- **Best Practice** 💡: Always **reset_index(drop=True)** before merges to avoid misalignment.



Select rows & columns – label vs position

```
Pandas

df = pd.DataFrame({'ID': ['r1', 'r2', 'r3'],
                   'Name': ['Alice', 'Bob', 'Cara'],
                   'Score': [85, 92, 78]}).set_index('ID')

# label-based
row_r2 = df.loc['r2']
cols_NameScore = df.loc[:, ['Name', 'Score']]

# position-based
first_row = df.iloc[0]
subset = df.iloc[0:2, 0:2] # first 2 rows, first 2 cols
```

Result (example)

ID	Name	Score	ID
r1	Alice	85	r1
r2	Bob	92	r2
r3	Cara	78	r3

- **`df.loc['r2']`** → Name: Bob, Score: 92
- **`df.iloc[0]`** → Name: Alice, Score: 85

Explanation/Tip

- loc = label-based, iloc = position-based.
- Trick: **`.at`** and **`.iat`** are faster for single values.



Filtering / Boolean indexing (query, isin, between)

```
Pandas

df = pd.DataFrame({'EmpID': ['E01', 'E02', 'E03'],
                  'Age': [28, 42, 33],
                  'Dept': ['Sales', 'HR', 'Sales']})

# Age > 30
df_age_over30 = df[df['Age'] > 30]
# multiple conditions
df_sales_over30 = df[(df['Age'] > 30) & (df['Dept'] == 'Sales')]
# query
df.query("Age > 30 and Dept == 'Sales'")
# isin
df[df['Dept'].isin(['Sales', 'HR'])]
# between
df[df['Age'].between(30, 40)]
```

Naveen Dharwan

Result

- **df_age_over30** rows: E02 (42), E03 (33)
- **df_sales_over30** rows: E03
- **df.query("Age > 30 and Dept == 'Sales'")** → rows: E03 (33, Sales)
- **df[df['Dept'].isin(['Sales', 'HR'])]** → rows: E03 (33, Sales)
- **df[df['Age'].between(30, 40)]** → rows: E03 (33, Sales)

Tip

- **query()** is handy for chained conditions when column names are simple identifiers.
- Common Error : Forgetting parentheses when using multiple conditions → **&** and **|** require brackets.



Missing Values

```
Pandas

df = pd.DataFrame({'ID':[1,2,3],
                  'Price':[250, float('nan'), 400],
                  'Discount':[10,5, float('nan')]})

# detect
missing_counts = df.isnull().sum()

# fill numeric with mean
df['Price'].fillna(df['Price'].mean(), inplace=True)

# fill specific per column
df.fillna({'Discount':0}, inplace=True)

# drop rows with any NaN
df_dropped = df.dropna(axis=0, how='any')

# interpolate (if time series)
# df['some'] = df['some'].interpolate()
```

Naveen Dhawan

Result

- missing_counts → Price:1, Discount:1
- After fill: Price(2) becomes mean of [250,400]=325

Tip

- ☒ Use **.isnull().sum()** at the top of any cleaning cell.
- Use **interpolate()** for time-series data.
- Fill missing values in-place to save memory.



Duplicates & replace

```
Pandas

df = pd.DataFrame({'UserID': ['U1', 'U2', 'U1'],
                  'Email':
                  ['alice@mail.com', 'bob@mail.com', 'alice@mail.com'],
                  'Country': ['India', 'USA', 'India']})

# find duplicates
dups = df[df.duplicated()]

# drop duplicates (keep first)
df_unique = df.drop_duplicates()

# replace values
df['Country'].replace({'USA': 'United States'},
inplace=True)
```

Result (example)

- **df.duplicated()** → True at row3
(dups → third row flagged as duplicate)
- **df.drop_duplicates()** → rows: 1, 2
- **df.replace('USA','United States')** → rows updated
(After replace: 'USA' → 'United States')

Tip

- Use a **subset** in duplicates to consider only certain columns
- **.replace()** can efficiently map multiple values.
- ⚠ Always inspect **value_counts()** before replace().



Adding / Dropping columns and rows

```
Pandas

df = pd.DataFrame(
    {"Order": [1, 2, 3], "Item": ["Pens", "Pencils",
    "Eraser"], "Qty": [10, 5, 2]}
)

# add new column
df["Price"] = [1.5, 0.5, 0.2]

# insert at position
df.insert(2, "Category", ["Office", "Office", "Office"])

# add row (concat)
new_row = pd.DataFrame(
    [{"Order": 4, "Item": "Notebook", "Qty": 7, "Price": 2.0,
    "Category": "Office"}]
)
df = pd.concat([df, new_row], ignore_index=True)

# drop column
df2 = df.drop("Price", axis=1)
```

Naveen Dhawan

Result (example)

- **df_concat** → rows: 1 (Pens, Office, 10, 1.5), 2 (Pencils, Office, 5, 0.5), 3 (Eraser, Office, 2, 0.2), 4 (Notebook, Office, 7, 2.0)
- New Price column added; concat created a 4th row.

Tip

- **append()** is deprecated. Instead, use **pd.concat([df, new_row])** for consistency and future-proof code.
- Quick before/after table showing assign() result.
- insert() can maintain column order if needed.



GroupBy basics & aggregation

```
Pandas

df = pd.DataFrame({'Country': ['India', 'India', 'USA'],
                   'Product': ['Phone', 'Laptop', 'Phone'],
                   'Sales': [200, 400, 300]})

# total sales per country
by_country = df.groupby('Country')['Sales'].sum()

# multiple agg per product
agg_prod = df.groupby('Product').agg({'Sales':
['sum', 'mean', 'count']})

# transform: broadcast mean per country
df['country_mean'] = df.groupby('Country')
['Sales'].transform('mean')
```

Naren Dhanan

Result

`df.groupby('Country')['Sales'].sum()` → India:600, USA:300

`df.groupby('Product').agg({'Sales': ['sum', 'mean', 'count']})` → Sales
sum/mean/count per product

Tip

- Use **agg()** for multiple aggregations in one step.
- Use **transform()** to broadcast aggregation results to original rows.



Merge / Join / Concat

```
Pandas

df1 = pd.DataFrame({'CustID': ['C1', 'C2', 'C3'],
                    'Name': ['Alice', 'Bob', 'Cara'],
                    'City': ['Pune', 'Delhi', 'Mumbai']})
df2 = pd.DataFrame({'CustID': ['C1', 'C2', 'C4'],
                    'OrderID': ['O1001', 'O1002', 'O1004'],
                    'Amount': [250, 150, 300]})

# inner join
inner = pd.merge(df1, df2, on='CustID', how='inner')

# left join
left = pd.merge(df1, df2, on='CustID', how='left')

# outer with indicator
outer = pd.merge(df1, df2, on='CustID', how='outer',
indicator=True)
```

Result (example)

CustID	Name	City	OrderID	Amount
C1	Alice	Pune	O1001	250
C2	Bob	Delhi	O1002	150

df_merge → rows: C1 (Alice, Pune, O1001, 250), C2 (Bob, Delhi, O1002, 150)

Tip

- Use **indicator=True** to see the merge origin; use suffixes when overlapping column names exist.



Pivot table

```
Pandas

df = pd.DataFrame({'Region': ['East', 'East', 'West'],
                    'Month': ['Jan', 'Feb', 'Jan'],
                    'Sales': [100, 120, 150]})

# pivot table
pivot = df.pivot_table(values='Sales', index='Region',
                        columns='Month', aggfunc='sum', fill_value=0)

# melt (wide -> long) example: assume pivoted data exists
melted = pd.melt(pivot.reset_index(), id_vars=
['Region'], value_vars=['Jan', 'Feb'], var_name='Month',
value_name='Sales')
```

Result

Region	Jan	Feb
East	100	120
West	150	0

pivot_table → East: (Jan=100, Feb=120), West: (Jan=150, Feb=0)

Tip

- Use pivot_table with aggfunc to handle duplicate index/column pairs.
- Use fill_value=0 in pivot_table to avoid NaNs.
- Use pivot_table to analyse monthly sales performance by region (great for sales dashboards).



Rolling / Shift / Diff

```
Pandas

df = pd.DataFrame({'Date': ['2023-01-01', '2023-01-02', '2023-01-03'],
                   'Sales': [100, 150, 120]})

df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)

df['Sales_rolling3'] = df['Sales'].rolling(window=3).mean()
df['Sales_shift1'] = df['Sales'].shift(1)
df['Sales_diff'] = df['Sales'].diff()
```

Result (example)

Date	Sales	Sales_rolling3	Sales_shift1	Sales_diff
01/01/2023	100	NaN	NaN	NaN
02/01/2023	150	NaN	100	50
03/01/2023	120	123.33	150	-30

Tip

- Set ***min_periods*** to control NaN at start; use *shift()* for lag features in ML
- Useful for moving averages, lagged features, and percent changes.



String Methods

```
Pandas

df = pd.DataFrame({'FullName': ['Alice Johnson', 'Bob S.', 'Charlie Brown'],
                  'Email': ['Alice@example.com', 'bob.smith@mail.com', 'charlie@site.org'],
                  'Note': ['VIP', 'new', 'loyal']})

# lower / strip / len
df['email_clean'] = df['Email'].str.lower().str.strip()
df['first_name'] = df['FullName'].str.split().str[0]
df['has_vip'] = df['Note'].str.contains('VIP', na=False)
# regex extract (domain)
df['domain'] = df['Email'].str.extract(r'@(.+)\$')
```

Nareen Dharwan

Result

FullName	email_clean	first_name	has_vip	domain
Alice Johnson	alice@example.com	Alice	TRUE	example.com

Tip

- **.str** methods are vectorised (fast) — avoid Python loops.
- Avoid Python loops over strings; slow for large datasets.



DateTime

```
Pandas

df = pd.DataFrame({'Date': ['2023-01-01', '2023-01-02', '2023-01-03'],
                   'Orders': [10, 12, 8],
                   'Revenue': [1000, 1200, 800]})

# parse dates
df['Date'] = pd.to_datetime(df['Date'])

# extract
df['Year'] = df['Date'].dt.year
df['DayName'] = df['Date'].dt.day_name()

# set index and resample (example)
df.set_index('Date', inplace=True)
monthly = df.resample('M').sum()
```

Naveen Dhawan

Result (example)

df['Year'] → 2023 for all rows

df['DayName'] → weekday names of given dates

df.resample('M').sum() → Orders=30, Revenue=3000

monthly will be aggregated (for the given sample, all in Jan → same numbers)

Tip

- Use **df['col'].dt** for fast extraction of date components.
- Parse dates at read time (parse_dates=['col']).
- Use pd.to_datetime(..., errors='coerce') for safe conversion.



apply / applymap / vectorized functions

```
Pandas

df = pd.DataFrame({'A':[1,4,7], 'B':[2,5,8], 'C':[3,6,9]})

# apply column-wise (sum of row)
df['row_sum'] = df.apply(lambda row:
row['A']+row['B']+row['C'], axis=1)

# applymap elementwise
df_str = df[['A', 'B', 'C']].applymap(lambda x: f'v{x}')

# vectorized operation (fast)
df['A_plus_1'] = df['A'] + 1
```

Naveen Dharwan

Result


row_sum values: 6, 15, 24

df.applymap(str) → all values converted to string

df.apply(lambda x:x+1) → A_plus_1 = [2,5,8]

Tip

- Prefer vectorized operations; apply is slower but useful for custom row-wise logic.
- Use applymap for element-wise operations on entire DataFrame.



Mapping / Encoding

```
Pandas

df = pd.DataFrame({'Person': ['A', 'B', 'C'],
                  'Gender': ['M', 'F', 'M'],
                  'Color': ['Red', 'Blue', 'Green']})

# map
df['Gender_code'] = df['Gender'].map({'M':0, 'F':1})

# factorize
df['Color_code'], uniques = pd.factorize(df['Color'])

# get_dummies
df_dummies = pd.get_dummies(df, columns=['Color'],
                             prefix='clr')
```

Result (example)

Person	Gender	Gender_code	Color_code
A	M	0	0

df['Gender_code'] → Male=0, Female=1

pd.factorize(Color) → Blue=0, Green=1, Red=2

pd.get_dummies → clr_Blue, clr_Green, clr_Red one-hot

Tip

- Use **map()** or **factorize()** for categorical encoding.
- **get_dummies()** for one-hot encoding; drop one column to avoid multicollinearity.



Outlier Detection (IQR) & filtering

```
Pandas

import pandas as pd

df = pd.DataFrame({'ID':[1,2,3], 'Value':
[10,1000,12]})

Q1, Q3 = df['Value'].quantile([0.25, 0.75])
IQR = Q3 - Q1

outliers = df[(df['Value'] < Q1 - 1.5*IQR) |
(df['Value'] > Q3 + 1.5*IQR)]
print("Q1:", Q1)
print("Q3:", Q3)
print("IQR:", IQR)
print("\nOutliers:\n", outliers)
```

Result (example)

Q1: 11.0
Q3: 505.0
IQR: 494.0

Outliers:
ID Value
1 2 1000



Q1 = 11.0, Q3 = 505.0, IQR = 494.0
The value 1000 is outside the
upper bound ($Q3 + 1.5 \times IQR = 1246.0$), hence detected as an
outlier.

Tip

- Use robust methods for skewed distributions (log transform before IQR if needed).
- Use IQR or z-score to detect outliers.
- **Best Practices** 💡: Visualise with boxplots before filtering.



MultIndex & hierarchical indexing

```
Pandas

df = pd.DataFrame({'Country': ['India', 'India', 'USA'],
                  'City': ['Pune', 'Mumbai', 'NY'],
                  'Sales': [100, 200, 150]})

df_multi = df.set_index(['Country', 'City'])
# xs cross-section
pune = df_multi.xs(('India', 'Pune'))
# swap levels
swapped = df_multi.swaplevel(0,1).sort_index()
```

Naveen Dhawan

Result

outliers → row with Value=1000

df_multi.loc[('India','Pune'),'Sales'] → 100

Tip

- MultiIndex is great for pivoted/aggregated results
- Use .xs() to select cross-sections.
- Re-index after explode and drop empty strings first.
- Use str.split(',', expand=True) to split multi-valued strings to columns.



Reshape: explode, split & expand

```
Pandas

df = pd.DataFrame({'ID':[1,2,3],
                  'Tags':[['A','B'], ['B','C','D'], ['A']],
                  'Notes':['ok','check','done']})

# explode
df_exploded = df.explode('Tags').reset_index(drop=True)

# split strings example
df_names = pd.DataFrame({'FullName':['Alice Johnson','Bob Smith','Cara Lee']})
df_names[['First','Last']] =
df_names['FullName'].str.split(' ', n=1, expand=True)
```

Naveen Dhanwan

Result (example)

ID	Tags	Notes
1	A	ok
1	B	ok
2	B	check

Tip

- explode() is ideal for list-like cells.
- After explode, re-index and cleanup may be needed for joins/groupby



Sorting & ranking

```
Pandas

df = pd.DataFrame({'Item': ['P1', 'P2', 'P3'],
                   'Sales': [200, 500, 150],
                   'Rating': [4.5, 4.9, 3.9]})

# sort
df_sorted = df.sort_values(['Sales'], ascending=False)

# rank
df['Sales_rank'] = df['Sales'].rank(ascending=False)

# top 2
top2 = df.nlargest(2, 'Sales')
```

Naveen Dhawan

Result

df.sort_values('Sales', ascending=False) → P2(500), P1(200), P3(150)

df.nlargest(2, 'Sales') → P2(500), P1(200)


df['Sales_rank'] → P2=1, P1=2, P3=3

Tip

- Use nlargest / nsmallest for top-k selection; faster than sort.
- .rank() useful for percentile ranking and relative metrics.



Memory & performance



```
df = pd.DataFrame({'ID':[1,2,3],
                  'Category':['A','B','A'],
                  'Value':[100,200,150]})

# downcast numeric
df['ID'] = pd.to_numeric(df['ID'], downcast='integer')
df['Value'] = pd.to_numeric(df['Value'],
                           downcast='unsigned')

# category dtype
df['Category'] = df['Category'].astype('category')
```

Naveen Dhanwan

Result (example)

df.info() shows smaller memory usage; Category stored efficiently.

ID dtype → int16

Value dtype → uint8

Category dtype → category

Tip

- Drop unused columns early to save memory.
- Downcast numeric types with `pd.to_numeric(..., downcast=...)`.



IO examples

```
Pandas

df = pd.DataFrame({'SKU': ['S1', 'S2', 'S3'],
                   'Price': [10.5, 5.0, 7.25],
                   'Stock': [100, 200, 50]})

# save
df.to_csv('inventory.csv', index=False)
df.to_excel('inventory.xlsx', index=False)
df.to_parquet('inventory.parquet', index=False)
```


Naveen Dhawan

Result

Files created → inventory.csv, inventory.xlsx, inventory.parquet
parquet is efficient for analytics pipelines.

Tip

- Always use index=False for portability unless index is meaningful
- Use chunksize for reading very large files.
- Use parquet for analytics pipelines; faster read/write than CSV.
- Test small samples before reading/writing large datasets.



Visualization (pandas + seaborn)

```
Pandas

import matplotlib.pyplot as plt
import seaborn as sns

df = pd.DataFrame({'Date': ['2023-01-01', '2023-01-02', '2023-01-03'],
                  'Sales': [100, 150, 120],
                  'Category': ['A', 'B', 'A']})
df['Date'] = pd.to_datetime(df['Date'])
# time series
df.set_index('Date')['Sales'].plot(title='Sales over Time')

# bar plot
df.groupby('Category')['Sales'].sum().plot(kind='bar')

# heatmap of correlation
sns.heatmap(df[['Sales']].corr(), annot=True)
plt.tight_layout()
```

Result (example)

Plots rendered → line (Sales over Time), bar (Sales by Category), heatmap (correlation)

Tip

- Plot only the necessary columns to improve performance.
- Use `plt.savefig('plot.png', dpi=300)` for LinkedIn images.
- Use `plt.tight_layout()` to avoid overlapping labels.



eval / query / pipe

```
Pandas

df = pd.DataFrame({'Qty':[2,1,3],
                  'Price':[100,200,50],
                  'Discount':[5,0,10]})

# eval to compute total
df.eval("Total = Qty * Price * (1 - Discount/100)", inplace=True)

# query
df.query("Total > 100")

# pipe example (functional)
def add_margin(df, pct):
    df['Price_margin'] = df['Price'] * (1 + pct)
    return df

df = df.pipe(add_margin, 0.1)
```

Naveen Dhawan

Result

Total computed via eval quickly
Total → [190, 200, 135]
Price_margin → [110, 220, 55]
df.query("Total > 100") → all rows

Tip

- eval & query can be faster for large frames; be mindful of names & namespaces.
- pipe() improves readability in method chains.



Crosstab & pivot examples

```
Pandas

df = pd.DataFrame({'User': ['U1', 'U2', 'U1'],
                    'Action': ['Click', 'View', 'View'],
                    'Count': [1, 2, 1]})

# crosstab
ct = pd.crosstab(df['User'], df['Action'], margins=True)

# pivot_table (sum)
pt = df.pivot_table(values='Count', index='User',
                     columns='Action', aggfunc='sum', fill_value=0)
```

Naveen Dharwan

Result (example)

Action	Click	View	All
U1	1	1	2
U2	0	2	2
All	1	3	4

Tip

- Use crosstab for frequency tables.
- Use margins=True for totals.
- Combine with normalize=True for percentages.



explode + list column analytics

```
Pandas

df = pd.DataFrame({'ArticleID': ['A1', 'A2', 'A3'],
                   'Tags': [ ['news', 'ai'],
                             ['ai', 'ml'], ['news', 'economy'] ],
                   'Title': ['AI today', 'ML basics', 'Market update']})

# explode tags to count frequencies
tags = df.explode('Tags')
tag_counts = tags['Tags'].value_counts()
```

Naveen Dhawan

Result

tag_counts → ai:2, news:2, ml:1, economy:1

Tip

- After explode, use groupby / value_counts to analyze lists.
- Useful for tag clouds, multiple-category analysis.
- explode() converts list-like cells into rows for aggregation



Error handling, assertions, and validation

```
Pandas

df = pd.DataFrame({'ID':[1,2,3], 'Value':[10,20,30]})

# check presence of column
assert 'Value' in df.columns, "Value column missing"

# check duplicates
if df['ID'].duplicated().any():
    raise ValueError('Duplicate IDs found')

# safe numeric conversion
df['Value'] = pd.to_numeric(df['Value'],
errors='raise')
```

Naveen Dharwan

Result (example)

Assertion passed, No duplicates, Values numeric converted successfully

Tip

- Use errors='raise' while debugging to fail fast; switch to coerce for graceful handling later.
- Use assert for sanity checks in ETL pipelines.



Useful one-liners

```
Pandas

df[['A', 'B']].applymap(str)
# elementwise transform
df.nlargest(3, 'Sales')
# top 3
df['col'].astype('category')
# memory & speed
pd.crosstab(df.A, df.B)
# contingency
df.sample(frac=0.1, random_state=42)
# random sample
```

Naveen Dhawan

Result

`applymap(str)` → all values strings
`nlargest(3,'Sales')` → top-3 Sales rows
`astype('category')` → dtype category
`pd.crosstab(df.A,df.B)` → contingency table
`df.sample(frac=0.1,random_state=42)` → 10% sample

Tip

- Quick data checks: `df.nunique()`, `df.value_counts()`.
- Memory check: `df.memory_usage(deep=True)`.
- Save a snippet file with these for quick copy/paste.



Final Summary & Quick Reference

Topic	Command Example	Notes
Read/Write	pd.read_csv / df.to_csv	use index=False for csv output
Select	df.loc / df.iloc	label vs integer
Missing	df.isnull / df.fillna	be careful with types

Key Commands

- Create: `pd.DataFrame()`
- Read: `pd.read_csv()`, `pd.read_excel()`
- Inspect: `df.head()`, `df.info()`, `df.describe()`
- Filter: `df[df.col > x]`, `df.query()`
- Group: `df.groupby(...).agg(...)`
- Merge: `pd.merge()`
- Reshape: `df.pivot_table()`, `pd.melt()`
- Time: `pd.to_datetime()`, `df.resample()`
- Save: `df.to_csv()`, `df.to_parquet()`

These tips are compiled from real-world projects where I used Pandas on 3,00,000+ row datasets.
Tested, optimised, and practical.

Save this cheat sheet for your next project